

hgb_features_modeling

September 21, 2025

1 Arrest Prediction — v1 (HistGradientBoosting + Engineered Features)

Goal: Beat RF v0 (PR-AUC 0.623) by adding time features, rare bucketing, and frequency encodings, then training a HistGradientBoosting baseline.

Dataset: data/processed/arrest_features.csv

Target: arrest (0/1)

Artifacts: saved to notebooks/artifacts/

```
[20]: # Core imports
import os, time, json, numpy as np, pandas as pd
from pathlib import Path

# Modeling + metrics
from sklearn.model_selection import train_test_split, StratifiedKFold, RandomizedSearchCV
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    average_precision_score, roc_auc_score, classification_report,
    confusion_matrix, precision_recall_curve, roc_curve
)
from sklearn.experimental import enable_hist_gradient_boosting
from sklearn.ensemble import HistGradientBoostingClassifier

import matplotlib.pyplot as plt
from scipy.stats import loguniform, randint
import tempfile
from sklearn.metrics import precision_recall_fscore_support

# Paths
REPO = Path.cwd()
while REPO.name != "chicago-crime-pipeline" and REPO.parent != REPO:
    REPO = REPO.parent
DATA = REPO / "data" / "processed"
ART = REPO / "notebooks" / "artifacts"
```

```

ART.mkdir(parents=True, exist_ok=True)

# Load
df = pd.read_csv(DATA / "arrest_features.csv")
assert "arrest" in df.columns
print(df.shape, df["arrest"].value_counts(dropna=False).to_dict())

# Split (same seed/stratify as v0)
TARGET = "arrest"
y = df[TARGET].astype(int).values
X = df.drop(columns=[TARGET]).copy()
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
X_train.shape, X_test.shape

```

(10482, 10) {0: 8972, 1: 1510}

[20]: ((8385, 9), (2097, 9))

```

[21]: def slice_metrics(X_df, y_true, proba, threshold, slice_col, min_support=40):
    """
    Compute precision/recall/F1 per value of a categorical slice column.
    Saves nothing; just returns a DataFrame. You can write it to CSV after.
    """
    if slice_col not in X_df.columns:
        print(f"[skip] slice column not found: {slice_col}")
        return None

    df = pd.DataFrame({
        slice_col: X_df[slice_col],
        "y": y_true,
        "pred": (proba >= threshold).astype(int)
    })

    rows = []
    for val, g in df.groupby(slice_col):
        n = len(g)
        if n < min_support:
            continue
        p, r, f1, _ = precision_recall_fscore_support(
            g["y"], g["pred"], average="binary", zero_division=0
        )
        rows.append({
            slice_col: val, "support": int(n),
            "precision": float(p), "recall": float(r), "f1": float(f1)
        })

```

```

if not rows:
    print(f"[note] no slices with support {min_support} for {slice_col}")
    return None

return pd.DataFrame(rows).sort_values("f1", ascending=False).
↪reset_index(drop=True)

```

```
[ ]:
```

Saved slice metrics for weekday → /Volumes/easystore/Projects/chicago-crime-pipeline/notebooks/artifacts/slice_metrics_weekday_hgb_v1_2025090921-232438.csv
 Saved slice metrics for hour_bin → /Volumes/easystore/Projects/chicago-crime-pipeline/notebooks/artifacts/slice_metrics_hour_bin_hgb_v1_2025090921-232438.csv
 Saved slice metrics for primary_type → /Volumes/easystore/Projects/chicago-crime-pipeline/notebooks/artifacts/slice_metrics_primary_type_hgb_v1_2025090921-232438.csv

=== weekday: top 5 by F1 ===

/var/folders/6z/l9wv3crd4n5bzgcrdd7vxq8m0000gn/T/ipykernel_79209/847869267.py:17
 : FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
for val, g in df.groupby(slice_col):
```

| | weekday | support | precision | recall | f1 |
|---|----------|---------|-----------|----------|----------|
| 0 | Saturday | 259 | 0.807692 | 0.700000 | 0.750000 |
| 1 | Friday | 311 | 0.682927 | 0.651163 | 0.666667 |
| 2 | Thursday | 379 | 0.701754 | 0.634921 | 0.666667 |
| 3 | Sunday | 298 | 0.727273 | 0.558140 | 0.631579 |
| 4 | Monday | 288 | 0.625000 | 0.571429 | 0.597015 |

=== weekday: bottom 5 by F1 (support 40) ===

| | weekday | support | precision | recall | f1 |
|---|-----------|---------|-----------|----------|----------|
| 2 | Thursday | 379 | 0.701754 | 0.634921 | 0.666667 |
| 3 | Sunday | 298 | 0.727273 | 0.558140 | 0.631579 |
| 4 | Monday | 288 | 0.625000 | 0.571429 | 0.597015 |
| 5 | Wednesday | 286 | 0.615385 | 0.533333 | 0.571429 |
| 6 | Tuesday | 276 | 0.490196 | 0.581395 | 0.531915 |

=== hour_bin: top 5 by F1 ===

| | hour_bin | support | precision | recall | f1 |
|---|----------|---------|-----------|----------|----------|
| 0 | 18-23 | 622 | 0.666667 | 0.626263 | 0.645833 |
| 1 | 06-11 | 426 | 0.666667 | 0.603774 | 0.633663 |
| 2 | 12-17 | 638 | 0.626374 | 0.612903 | 0.619565 |
| 3 | 00-05 | 411 | 0.659574 | 0.543860 | 0.596154 |

=== hour_bin: bottom 5 by F1 (support 40) ===

| | hour_bin | support | precision | recall | f1 |
|---|----------|---------|-----------|----------|----------|
| 0 | 18-23 | 622 | 0.666667 | 0.626263 | 0.645833 |
| 1 | 06-11 | 426 | 0.666667 | 0.603774 | 0.633663 |
| 2 | 12-17 | 638 | 0.626374 | 0.612903 | 0.619565 |
| 3 | 00-05 | 411 | 0.659574 | 0.543860 | 0.596154 |

=== primary_type: top 5 by F1 ===

| | primary_type | support | precision | recall | f1 |
|---|-------------------|---------|-----------|----------|----------|
| 0 | NARCOTICS | 58 | 0.913793 | 1.000000 | 0.954955 |
| 1 | WEAPONS VIOLATION | 53 | 0.730769 | 1.000000 | 0.844444 |
| 2 | ROBBERY | 56 | 0.800000 | 0.571429 | 0.666667 |
| 3 | OTHER OFFENSE | 134 | 0.515152 | 0.680000 | 0.586207 |
| 4 | CRIMINAL TRESPASS | 43 | 0.392857 | 0.916667 | 0.550000 |

=== primary_type: bottom 5 by F1 (support 40) ===

| | primary_type | support | precision | recall | f1 |
|----|---------------------|---------|-----------|----------|----------|
| 7 | BATTERY | 379 | 0.607143 | 0.283333 | 0.386364 |
| 8 | THEFT | 473 | 0.419355 | 0.325000 | 0.366197 |
| 9 | CRIMINAL DAMAGE | 232 | 1.000000 | 0.111111 | 0.200000 |
| 10 | ASSAULT | 187 | 0.400000 | 0.090909 | 0.148148 |
| 11 | MOTOR VEHICLE THEFT | 160 | 0.000000 | 0.000000 | 0.000000 |

```
[31]: # Rebuild engineered features from X_train/X_test
X_train_fe = X_train.copy()
X_test_fe = X_test.copy()

# Weekday
for Xdf in (X_train_fe, X_test_fe):
    Xdf["weekday"] = pd.to_datetime(Xdf["date"]).dt.day_name()

# Hour bins (ensure object dtype)
bins = [0,6,12,18,24]
labels = ["00-05", "06-11", "12-17", "18-23"]
for Xdf in (X_train_fe, X_test_fe):
    Xdf["hour_bin"] = pd.cut(Xdf["hour"].astype(int), bins=bins, right=False,
↪ labels=labels).astype(object)

# Rare bucket helper
def rare_bucket(train_col, test_col, min_count=40):
    vc = train_col.value_counts()
    keep = set(vc[vc >= min_count].index)
    return (train_col.where(train_col.isin(keep), "__RARE__"),
            test_col.where(test_col.isin(keep), "__RARE__"))

# Rare bucket base categoricals if present
for col in ["location_description", "primary_type"]:
```

```

    if col in X_train_fe.columns:
        X_train_fe[col], X_test_fe[col] = rare_bucket(X_train_fe[col],
↪X_test_fe[col], 40)

# Frequency encodes
def add_freq_encode(col):
    freq = X_train_fe[col].astype(object).value_counts(normalize=True)
    X_train_fe[f"{col}_freq"] = X_train_fe[col].map(freq).astype("float64").
↪fillna(0.0).to_numpy()
    X_test_fe[f"{col}_freq"] = X_test_fe[col].map(freq).astype("float64").
↪fillna(0.0).to_numpy()

for col in ["primary_type", "location_description", "weekday", "hour_bin"]:
    if col in X_train_fe.columns: add_freq_encode(col)

# Target mean for primary_type
if "primary_type" in X_train_fe.columns:
    arrest_rate = pd.Series(y_train).groupby(X_train_fe["primary_type"]).mean()
    X_train_fe["ptype_arrest_rate"] = X_train_fe["primary_type"].
↪map(arrest_rate)
    X_test_fe["ptype_arrest_rate"] = X_test_fe["primary_type"].
↪map(arrest_rate).fillna(float(arrest_rate.mean()))
else:
    X_train_fe["ptype_arrest_rate"] = 0.0
    X_test_fe["ptype_arrest_rate"] = 0.0

# Interaction primary_type × hour_bin (rare-bucket)
if set(["primary_type", "hour_bin"]).issubset(X_train_fe.columns):
    X_train_fe["ptype_x_hourbin"] = X_train_fe["primary_type"].astype(str) +
↪"_" + X_train_fe["hour_bin"].astype(str)
    X_test_fe["ptype_x_hourbin"] = X_test_fe["primary_type"].astype(str) +
↪"_" + X_test_fe["hour_bin"].astype(str)
    X_train_fe["ptype_x_hourbin"], X_test_fe["ptype_x_hourbin"] = rare_bucket(
        X_train_fe["ptype_x_hourbin"], X_test_fe["ptype_x_hourbin"], 30
    )
else:
    X_train_fe["ptype_x_hourbin"] = "__MISSING__"
    X_test_fe["ptype_x_hourbin"] = "__MISSING__"

```

```

[32]: # Your lists
cat_cols_fe =
↪["date", "primary_type", "location_description", "location_grouped", "weekday", "hour_bin", "ptype_
num_cols_fe =
↪["id", "year", "month", "dow", "hour", "primary_type_freq", "location_description_freq", "weekday_

# Remove columns that don't exist (e.g., location_grouped may be absent)

```

```

present = set(X_train_fe.columns)
cat_cols_used = [c for c in cat_cols_fe if c in present]
num_cols_used = [c for c in num_cols_fe if c in present]

print("Using categorical:", cat_cols_used)
print("Using numeric      :", num_cols_used)

pre_fe = ColumnTransformer(
    transformers=[
        ("cat", OneHotEncoder(handle_unknown="ignore", sparse_output=False),
        ↪cat_cols_used),
        ("num", "passthrough", num_cols_used),
    ],
    remainder="drop",
    verbose_feature_names_out=False,
)

```

Using categorical: ['date', 'primary_type', 'location_description', 'location_grouped', 'weekday', 'hour_bin', 'ptype_x_hourbin']
Using numeric : ['id', 'year', 'month', 'dow', 'hour', 'primary_type_freq', 'location_description_freq', 'weekday_freq', 'hour_bin_freq', 'ptype_arrest_rate']

```

[30]: from sklearn.model_selection import train_test_split

# Subsample ~5k rows for faster search
SUB_N = 5000
if len(y_train) > SUB_N:
    X_sub, _, y_sub, _ = train_test_split(
        X_train_fe, y_train, train_size=SUB_N,
        stratify=y_train, random_state=42
    )
else:
    X_sub, y_sub = X_train_fe, y_train

hgb_search = RandomizedSearchCV(
    hgb_pipe, param_distributions=param_dist,
    n_iter=6,                # 6 candidates
    scoring="average_precision",
    refit=True, cv=2,        # 2 folds
    n_jobs=-1, random_state=42, verbose=2
)

hgb_search.fit(X_sub, y_sub, clf__sample_weight=sw_train[:len(y_sub)])
print("Best HGB params:", hgb_search.best_params_)
print("Best CV PR-AUC:", round(hgb_search.best_score_, 4))

```

Fitting 2 folds for each of 6 candidates, totalling 12 fits

```

[CV] END clf__l2_regularization=0.0003487351559952693,
clf__learning_rate=0.06207090305742937, clf__max_depth=5, clf__max_iter=90,
clf__max_leaf_nodes=39, clf__min_samples_leaf=176; total time= 0.0s
[CV] END clf__l2_regularization=0.00010062545641808922,
clf__learning_rate=0.1978522015446167, clf__max_depth=3, clf__max_iter=91,
clf__max_leaf_nodes=41, clf__min_samples_leaf=81; total time= 0.0s
[CV] END clf__l2_regularization=0.00010062545641808922,
clf__learning_rate=0.1978522015446167, clf__max_depth=3, clf__max_iter=91,
clf__max_leaf_nodes=41, clf__min_samples_leaf=81; total time= 0.0s
[CV] END clf__l2_regularization=0.0020059560245279666,
clf__learning_rate=0.18679147494991152, clf__max_depth=5, clf__max_iter=87,
clf__max_leaf_nodes=44, clf__min_samples_leaf=80; total time= 0.0s
[CV] END clf__l2_regularization=0.012306931514988042,
clf__learning_rate=0.1334357418214006, clf__max_depth=7, clf__max_iter=81,
clf__max_leaf_nodes=39, clf__min_samples_leaf=97; total time= 0.0s
[CV] END clf__l2_regularization=0.0003487351559952693,
clf__learning_rate=0.06207090305742937, clf__max_depth=5, clf__max_iter=90,
clf__max_leaf_nodes=39, clf__min_samples_leaf=176; total time= 0.0s
[CV] END clf__l2_regularization=0.012306931514988042,
clf__learning_rate=0.1334357418214006, clf__max_depth=7, clf__max_iter=81,
clf__max_leaf_nodes=39, clf__min_samples_leaf=97; total time= 0.0s
[CV] END clf__l2_regularization=0.0020059560245279666,
clf__learning_rate=0.18679147494991152, clf__max_depth=5, clf__max_iter=87,
clf__max_leaf_nodes=44, clf__min_samples_leaf=80; total time= 0.0s
[CV] END clf__l2_regularization=0.00010582064396389428,
clf__learning_rate=0.051624394129231366, clf__max_depth=5, clf__max_iter=138,
clf__max_leaf_nodes=46, clf__min_samples_leaf=74; total time= 0.0s
[CV] END clf__l2_regularization=0.00010582064396389428,
clf__learning_rate=0.051624394129231366, clf__max_depth=5, clf__max_iter=138,
clf__max_leaf_nodes=46, clf__min_samples_leaf=74; total time= 0.0s
[CV] END clf__l2_regularization=0.003853103152262984,
clf__learning_rate=0.1484885740998741, clf__max_depth=5, clf__max_iter=123,
clf__max_leaf_nodes=38, clf__min_samples_leaf=123; total time= 0.0s
[CV] END clf__l2_regularization=0.003853103152262984,
clf__learning_rate=0.1484885740998741, clf__max_depth=5, clf__max_iter=123,
clf__max_leaf_nodes=38, clf__min_samples_leaf=123; total time= 0.0s

```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[30], line 21
     11 X_sub, y_sub = X_train_fe, y_train
     13 hgb_search = RandomizedSearchCV(
     14     hgb_pipe, param_distributions=param_dist,
     15     n_iter=6,                                # 6 candidates
    (...) 18     n_jobs=-1, random_state=42, verbose=2
     19 )
--> 21 hgb_search.fit(X_sub, y_sub, clf__sample_weight=sw_train[:len(y_sub)])

```

```

22 print("Best HGB params:", hgb_search.best_params_)
23 print("Best CV PR-AUC:", round(hgb_search.best_score_, 4))

File /Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/
↳site-packages/sklearn/base.py:1365, in _fit_context.<locals>.decorator.
↳<locals>.wrapper(estimator, *args, **kwargs)
1358     estimator._validate_params()
1360 with config_context(
1361     skip_parameter_validation=(
1362         prefer_skip_nested_validation or global_skip_validation
1363     )
1364 ):
-> 1365     return fit_method(estimator, *args, **kwargs)

File /Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/
↳site-packages/sklearn/model_selection/_search.py:1051, in BaseSearchCV.
↳fit(self, X, y, **params)
1045     results = self._format_results(
1046         all_candidate_params, n_splits, all_out, all_more_results
1047     )
1049     return results
-> 1051 self._run_search(evaluate_candidates)
1053 # multimetric is determined here because in the case of a callable
1054 # self.scoring the return type is only known after calling
1055 first_test_score = all_out[0]["test_scores"]

File /Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/
↳site-packages/sklearn/model_selection/_search.py:1992, in RandomizedSearchCV.
↳_run_search(self, evaluate_candidates)
1990 def _run_search(self, evaluate_candidates):
1991     """Search n_iter candidates from param_distributions"""
-> 1992     evaluate_candidates(
1993         ParameterSampler(
1994             self.param_distributions, self.n_iter, random_state=self.random_state
1995         )
1996     )

File /Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/
↳site-packages/sklearn/model_selection/_search.py:1028, in BaseSearchCV.fit.
↳<locals>.evaluate_candidates(candidate_params, cv, more_results)
1021 elif len(out) != n_candidates * n_splits:
1022     raise ValueError(
1023         "cv.split and cv.get_n_splits returned "
1024         "inconsistent results. Expected {} "
1025         "splits, got {}".format(n_splits, len(out) // n_candidates)
1026     )
-> 1028 warn_or_raise_about_fit_failures(out, self.error_score)
1030 # For callable self.scoring, the return type is only know after

```



```

1031 # calling. If the return type is a dictionary, the error scores
1032 # can now be inserted with the correct key. The type checking
1033 # of out will be done in `_insert_error_scores`.
1034 if callable(self.scoring):

```

```

File /Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/
↳ site-packages/sklearn/model_selection/_validation.py:505, in _
↳ warn_or_raise_about_fit_failures(results, error_score)
    498 if num_failed_fits == num_fits:
    499     all_fits_failed_message = (
    500         f"\nAll the {num_fits} fits failed.\n"
    501         "It is very likely that your model is misconfigured.\n"
    502         "You can try to debug the error by setting error_score='raise'.
↳ \n\n"
    503         f"Below are more details about the failures:
↳ \n{fit_errors_summary}"
    504     )
--> 505     raise ValueError(all_fits_failed_message)
    507 else:
    508     some_fits_failed_message = (
    509         f"\n{num_failed_fits} fits failed out of a total of {num_fits}.
↳ \n"
    510         "The score on these train-test partitions for these parameters"
    511         f"Below are more details about the failures:
↳ \n{fit_errors_summary}"
    512     )

```

ValueError:

All the 12 fits failed.

It is very likely that your model is misconfigured.

You can try to debug the error by setting error_score='raise'.

Below are more details about the failures:

12 fits failed with the following error:

Traceback (most recent call last):

```

File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↳ site-packages/pandas/core/indexes/base.py", line 3805, in get_loc
    return self._engine.get_loc(casted_key)
    ~~~~~

```

```

File "index.pyx", line 167, in pandas._libs.index.IndexEngine.get_loc
File "index.pyx", line 196, in pandas._libs.index.IndexEngine.get_loc
File "pandas/_libs/hashtable_class_helper.pxi", line 7081, in pandas._libs.
↳ hashtable.PyObjectHashTable.get_item
File "pandas/_libs/hashtable_class_helper.pxi", line 7089, in pandas._libs.
↳ hashtable.PyObjectHashTable.get_item

```

KeyError: 'ptype_arrest_rate'

The above exception was the direct cause of the following exception:

Traceback (most recent call last):

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/sklearn/utils/_indexing.py", line 443, in _get_column_indices
    col_idx = all_columns.get_loc(col)
    ~~~~~
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/pandas/core/indexes/base.py", line 3812, in get_loc
    raise KeyError(key) from err
```

KeyError: 'ptype_arrest_rate'

The above exception was the direct cause of the following exception:

Traceback (most recent call last):

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/sklearn/model_selection/_validation.py", line 859, in _
↪_fit_and_score
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/sklearn/base.py", line 1365, in wrapper
    return fit_method(estimator, *args, **kwargs)
    ~~~~~
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/sklearn/pipeline.py", line 655, in fit
    Xt = self._fit(X, y, routed_params, raw_params=params)
    ~~~~~
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/sklearn/pipeline.py", line 589, in _fit
    X, fitted_transformer = fit_transform_one_cached(
    ~~~~~
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/joblib/memory.py", line 607, in __call__
    return self._cached_call(args, kwargs, shelving=False)[0]
    ~~~~~
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/joblib/memory.py", line 562, in _cached_call
    return self._call(call_id, args, kwargs, shelving)
    ~~~~~
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/joblib/memory.py", line 832, in _call
    output = self.func(*args, **kwargs)
    ~~~~~
```

```
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↪site-packages/sklearn/pipeline.py", line 1540, in _fit_transform_one
    res = transformer.fit_transform(X, y, **params.get("fit_transform", {}))
    ~~~~~
```

```

File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↳site-packages/sklearn/utils/_set_output.py", line 316, in wrapped
    data_to_wrap = f(self, X, *args, **kwargs)
    ~~~~~
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↳site-packages/sklearn/base.py", line 1365, in wrapper
    return fit_method(estimator, *args, **kwargs)
    ~~~~~
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↳site-packages/sklearn/compose/_column_transformer.py", line 988, in _
↳fit_transform
    self._validate_column_callables(X)
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↳site-packages/sklearn/compose/_column_transformer.py", line 541, in _
↳_validate_column_callables
    transformer_to_input_indices[name] = _get_column_indices(X, columns)
    ~~~~~
File "/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12
↳site-packages/sklearn/utils/_indexing.py", line 451, in _get_column_indices
    raise ValueError("A given column is not a column of the dataframe") from e
ValueError: A given column is not a column of the dataframe

```

```

[13]: hgb_final = hgb_search.best_estimator_
hgb_final.fit(X_train_fe, y_train, clf__sample_weight=sw_train)

```

```

[13]: Pipeline(memory='/var/folders/6z/l9wv3crd4n5bzgcrdd7vxq8m0000gn/T/tmptrqagad5',
              steps=[('pre',
                      ColumnTransformer(transformers=[('cat',
OneHotEncoder(handle_unknown='ignore',
sparse_output=False),
                                                         ['date', 'primary_type',
'location_description',
'location_grouped',
'weekday', 'hour_bin']),
('num', 'passthrough',
['id', 'year', 'month', 'dow',
'hour', 'primary_type_freq',
'location_description_freq',
'weekday_freq',
'hour_bin_freq'])]),
                      verbose_feature_names_out=False)),
              ('clf',
               HistGradientBoostingClassifier(l2_regularization=np.float64(0.0
020059560245279666),
learning_rate=np.float64(0.18679147494991152),
max_depth=5, max_iter=87,

```

```
max_leaf_nodes=44,
min_samples_leaf=80,
random_state=42)))]
```

```
[14]: proba_hgb = hgb_final.predict_proba(X_test_fe)[:,-1]

print("HGB TEST PR-AUC:", round(average_precision_score(y_test, proba_hgb), 4))
print("HGB TEST ROC-AUC:", round(roc_auc_score(y_test, proba_hgb), 4))

# Threshold tuning
prec, rec, thr = precision_recall_curve(y_test, proba_hgb)
f1s = 2*prec*rec/(prec+rec+1e-12)
best_idx = np.nanargmax(f1s)
thr_hgb = thr[best_idx] if best_idx < len(thr) else 0.5
pred_hgb = (proba_hgb >= thr_hgb).astype(int)

print("Best threshold:", float(thr_hgb), "Best F1:", float(f1s[best_idx]))
print(classification_report(y_test, pred_hgb, digits=3))
print("Confusion:\n", confusion_matrix(y_test, pred_hgb))
```

HGB TEST PR-AUC: 0.6569

HGB TEST ROC-AUC: 0.8878

Best threshold: 0.6953415078096913 Best F1: 0.6265060240958864

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.934 | 0.946 | 0.940 | 1795 |
| 1 | 0.652 | 0.603 | 0.627 | 302 |
| accuracy | | | 0.897 | 2097 |
| macro avg | 0.793 | 0.774 | 0.783 | 2097 |
| weighted avg | 0.893 | 0.897 | 0.895 | 2097 |

Confusion:

```
[[1698  97]
 [ 120 182]]
```

```
[ ]: if 'ART' not in globals():
    ART = Path("notebooks/artifacts"); ART.mkdir(parents=True, exist_ok=True)
if 'stamp' not in globals():
    import time; stamp = time.strftime("%Y%m%d-%H%M%S")

cols_to_check = ["weekday", "hour_bin", "primary_type"] # adjust if needed
slice_tables = {}

for col in cols_to_check:
    tbl = slice_metrics(X_test_fe, y_test, proba_hgb, thr_hgb, col,
        ↪min_support=40)
```

```

if tbl is not None:
    slice_tables[col] = tbl
    out_path = ART / f"slice_metrics_{col}_hgb_v1_{stamp}.csv"
    tbl.to_csv(out_path, index=False)
    print(f"Saved slice metrics for {col} → {out_path}")

# (optional) quick peek in the notebook
for col, tbl in slice_tables.items():
    print(f"\n=== {col}: top 5 by F1 ===")
    display(tbl.head(5))
    print(f"=== {col}: bottom 5 by F1 (support 40) ===")
    display(tbl.tail(5))

```

```

[15]: stamp = time.strftime("%Y%m%d-%H%M%S")

metrics = {
    "timestamp": stamp,
    "model": "HGB + FE v1",
    "test_pr_auc": float(average_precision_score(y_test, proba_hgb)),
    "test_roc_auc": float(roc_auc_score(y_test, proba_hgb)),
    "threshold_tuned": float(thr_hgb),
    "confusion_tuned": confusion_matrix(y_test, pred_hgb).tolist(),
    "class_report_tuned": classification_report(y_test, pred_hgb,
    ↪ output_dict=True),
    "best_params": {k: (float(v) if hasattr(v, "item") else v) for k,v in
    ↪ hgb_search.best_params_.items()}
}

with open(ART / f"metrics_hgb_v1_{stamp}.json", "w") as f:
    json.dump(metrics, f, indent=2)

with open(ART / "decision_threshold_hgb_v1.txt", "w") as f:
    f.write(str(metrics["threshold_tuned"]))

# PR/ROC plots
prec, rec, _ = precision_recall_curve(y_test, proba_hgb)
fpr, tpr, _ = roc_curve(y_test, proba_hgb)

plt.figure(); plt.plot(rec, prec); plt.xlabel("Recall"); plt.ylabel("Precision")
plt.title(f"HGB PR curve (AP={metrics['test_pr_auc']:.3f})"); plt.grid(True,
    ↪ alpha=0.3)
plt.savefig(ART / f"pr_curve_hgb_v1_{stamp}.png", bbox_inches="tight"); plt.
    ↪ close()

plt.figure(); plt.plot(fpr, tpr); plt.plot([0,1],[0,1], '--')
plt.xlabel("FPR"); plt.ylabel("TPR"); plt.title(f"HGB ROC curve
    ↪ (AUC={metrics['test_roc_auc']:.3f})")

```

```
plt.grid(True, alpha=0.3)
plt.savefig(ART / f"roc_curve_hgb_v1_{stamp}.png", bbox_inches="tight"); plt.
    ↪close()
```

```
print("Saved HGB v1 artifacts:", ART)
```

Saved HGB v1 artifacts: /Volumes/easystore/Projects/chicago-crime-pipeline/notebooks/artifacts

```
[23]: for col in ["weekday", "hour_bin", "primary_type"]:
        out = slice_metrics(X_test_fe, y_test, proba_hgb, thr_hgb, col)
        if out is not None:
            out.to_csv(ART / f"slice_metrics_{col}_hgb_v1_{stamp}.csv", index=False)
            print(f"Saved slice metrics for {col}")
```

Saved slice metrics for weekday
 Saved slice metrics for hour_bin
 Saved slice metrics for primary_type

/var/folders/6z/l9wv3crd4n5bzgcrdd7vxq8m0000gn/T/ipykernel_79209/847869267.py:17
 : FutureWarning: The default of observed=False is deprecated and will be changed
 to True in a future version of pandas. Pass observed=False to retain current
 behavior or observed=True to adopt the future default and silence this warning.
 for val, g in df.groupby(slice_col):

```
[24]: def threshold_for_recall(y_true, proba, target=0.70):
        prec, rec, thr = precision_recall_curve(y_true, proba)
        idx = np.argmax(rec >= target)
        th = thr[max(idx-1, 0)] if idx < len(thr) else 0.5
        return float(th), float(prec[max(idx-1, 0)]), float(rec[max(idx-1, 0)])

        thr_r70, p_at_r70, r_at_r70 = threshold_for_recall(y_test, proba_hgb, target=0.
            ↪70)
        print("Threshold for recall 0.70:", thr_r70, "| precision", p_at_r70, "|
            ↪recall", r_at_r70)

        pred_r70 = (proba_hgb >= thr_r70).astype(int)
        print(classification_report(y_test, pred_r70, digits=3))
        print("Confusion:\n", confusion_matrix(y_test, pred_r70))
```

Threshold for recall 0.70: 0.0003309188924939405 | precision
 0.14401525989508823 | recall 1.0

| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| 0 | 0.000 | 0.000 | 0.000 | 1795 |
| 1 | 0.144 | 1.000 | 0.252 | 302 |
| accuracy | | | 0.144 | 2097 |
| macro avg | 0.072 | 0.500 | 0.126 | 2097 |

| | | | | |
|--------------|-------|-------|-------|------|
| weighted avg | 0.021 | 0.144 | 0.036 | 2097 |
|--------------|-------|-------|-------|------|

Confusion:

```
[[ 0 1795]
 [ 0 302]]
```

/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```

/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```

/Volumes/easystore/Projects/chicago-crime-pipeline/.venv/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
```

```
[25]: slice_tables_r70 = {}
      for col in slice_tables.keys():
          tbl = slice_metrics(X_test_fe, y_test, proba_hgb, thr_r70, col,
                               min_support=40)
          if tbl is not None:
              slice_tables_r70[col] = tbl
              display(tbl.head(5)); display(tbl.tail(5))
```

| | weekday | support | precision | recall | f1 |
|---|-----------|---------|-----------|--------|----------|
| 0 | Thursday | 379 | 0.166227 | 1.0 | 0.285068 |
| 1 | Wednesday | 286 | 0.157343 | 1.0 | 0.271903 |
| 2 | Tuesday | 276 | 0.155797 | 1.0 | 0.269592 |
| 3 | Sunday | 298 | 0.144295 | 1.0 | 0.252199 |
| 4 | Friday | 311 | 0.138264 | 1.0 | 0.242938 |

| | weekday | support | precision | recall | f1 |
|---|----------|---------|-----------|--------|----------|
| 2 | Tuesday | 276 | 0.155797 | 1.0 | 0.269592 |
| 3 | Sunday | 298 | 0.144295 | 1.0 | 0.252199 |
| 4 | Friday | 311 | 0.138264 | 1.0 | 0.242938 |
| 5 | Monday | 288 | 0.121528 | 1.0 | 0.216718 |
| 6 | Saturday | 259 | 0.115830 | 1.0 | 0.207612 |

/var/folders/6z/l9wv3crd4n5bzgcrdd7vxq8m0000gn/T/ipykernel_79209/847869267.py:17: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
for val, g in df.groupby(slice_col):
```

| | hour_bin | support | precision | recall | f1 |
|---|----------|---------|-----------|--------|----------|
| 0 | 18-23 | 622 | 0.159164 | 1.0 | 0.274619 |
| 1 | 12-17 | 638 | 0.145768 | 1.0 | 0.254446 |
| 2 | 00-05 | 411 | 0.138686 | 1.0 | 0.243590 |
| 3 | 06-11 | 426 | 0.124413 | 1.0 | 0.221294 |

| | hour_bin | support | precision | recall | f1 |
|---|----------|---------|-----------|--------|----------|
| 0 | 18-23 | 622 | 0.159164 | 1.0 | 0.274619 |
| 1 | 12-17 | 638 | 0.145768 | 1.0 | 0.254446 |
| 2 | 00-05 | 411 | 0.138686 | 1.0 | 0.243590 |
| 3 | 06-11 | 426 | 0.124413 | 1.0 | 0.221294 |

| | primary_type | support | precision | recall | f1 |
|---|-------------------|---------|-----------|--------|----------|
| 0 | NARCOTICS | 58 | 0.913793 | 1.0 | 0.954955 |
| 1 | WEAPONS VIOLATION | 53 | 0.716981 | 1.0 | 0.835165 |
| 2 | CRIMINAL TRESPASS | 43 | 0.279070 | 1.0 | 0.436364 |
| 3 | OTHER OFFENSE | 134 | 0.186567 | 1.0 | 0.314465 |
| 4 | BATTERY | 379 | 0.158311 | 1.0 | 0.273349 |

| | primary_type | support | precision | recall | f1 |
|----|---------------------|---------|-----------|--------|----------|
| 7 | THEFT | 473 | 0.084567 | 1.0 | 0.155945 |
| 8 | CRIMINAL DAMAGE | 232 | 0.038793 | 1.0 | 0.074689 |
| 9 | BURGLARY | 96 | 0.031250 | 1.0 | 0.060606 |
| 10 | DECEPTIVE PRACTICE | 136 | 0.029412 | 1.0 | 0.057143 |
| 11 | MOTOR VEHICLE THEFT | 160 | 0.018750 | 1.0 | 0.036810 |

```
[ ]: # target mean encoding for primary_type
arrest_rate = pd.Series(y_train).groupby(X_train_fe["primary_type"]).mean()
X_train_fe["ptype_arrest_rate"] = X_train_fe["primary_type"].map(arrest_rate)
X_test_fe["ptype_arrest_rate"] = X_test_fe["primary_type"].map(arrest_rate).
    ↪ fillna(arrest_rate.mean())

num_cols_fe.append("ptype_arrest_rate")

# Add interaction features to reduce time-of-day false positives\
# Combine hour_bin x primary_type
X_train_fe["ptype_x_hourbin"] = X_train_fe["primary_type"].astype(str) + "_" +
    ↪ X_train_fe["hour_bin"].astype(str)
X_test_fe["ptype_x_hourbin"] = X_test_fe["primary_type"].astype(str) + "_" +
    ↪ X_test_fe["hour_bin"].astype(str)
```

```
[ ]:
```